




Issue 1 — The timer issue

You ran out of time on consultations (the 10-minute timer expired before you could consult board members). The yellow warning *"Time has expired! Consultations are now locked"* appeared. Two problems:

- **Late-submission penalty isn't visible.** The text says "it will be recorded as a late submission" — but no indication of how that affects scoring
- **The candidate likely needs more time on Round 1** because they're still learning the simulation interface. Suggest making Round 1 either untimed or 15 minutes; tighten progressively in later rounds

Issue 2 — Mixed currency/scale warning is concerning

The left sidebar shows: "⚠️ *Mixed currency/scale units detected (₹, B, \$M, TB). Values are shown as extracted and are NOT normalized to a common scale.*"

This is a **data quality bug** that's getting exposed to the candidate. For a Module 6 simulation specifically, this matters — the candidate is supposed to be doing financial-ratio analysis, and they cannot if the units aren't consistent. If the dashboard mixes Indian Rupees (₹), Billions, Millions, and TB (data storage) without normalising, any ratio computation will be nonsensical. Tell your developer: **normalise all financial figures to a single unit (US\$ Millions recommended) before display.**

Issue 3 — "No high priority metrics flagged"

The High Priority Metrics panel says "No high priority metrics flagged." For a Round 1 prudence test, the high-priority metrics should be obvious:

- The \$75M confirmed liability
- The 40% stock drop
- 8% client churn (annualised)
- 1 major + 3 minor data security incidents

These are all flagged in the Current Challenges section but not surfaced as "high priority metrics." Either the simulation is missing the wiring to elevate them, or the concept of "high priority metric" isn't yet defined in the rubric. Worth asking your developer.

Issue 4 — The Module 6 vocabulary still isn't being scored explicitly

Look at Round 1's response from the agents: they use prudence-related language ("provisioning," "contingent liabilities," "recognition," "investor confidence") — but I cannot see anywhere in the UI that the simulation **scores** the candidate on whether they used:

- The phrase "Principle of Prudence" by name
- Reference to AS 5 disclosure standard
- Distinction between exceptional vs extraordinary items
- Recognition that Ind AS does not permit extraordinary classification

These are explicit M6 Session 1 testable concepts. The simulation is *implicitly* in the right territory but isn't *explicitly* scoring against M6 vocabulary. This matters for your developer: even if the candidate wins the debate, did they win it for *M6-correct reasons*?

Issue 5 — David Chen approval logic is missing from the deliberation

You're playing as David Chen. The deliberation shows 8 directors supporting, but **David Chen himself isn't listed in the support count** — which is correct (you can't vote on your own decision). But neither is he visible as the proposer/owner. The UI could be clearer about: "You proposed this decision. The board's response is below."

Issue 6 — The next round's M6 focus is still unknown

We won't know whether Rounds 2–5 cover Sessions 2, 3, 4 of M6 until you actually progress. The Round 1 mechanic is solid. The question is whether Round 2 says "*Focus: Revenue Recognition*" or "*Focus: Exceptional vs Extraordinary Items*" (good — M6 Session 2) or "*Focus: PR Strategy*" (bad — drifts off-module).

Add to Part A (Critical bugs):

A7. Currency normalisation bug. The dashboard shows mixed units (₹, B, M, TB) and surfaces a warning to the candidate that values are "NOT normalized to a common scale." For a Module 6 financial-ratios simulation specifically, this is fatal — ratios

cannot be computed across incompatible units. Normalise all monetary figures to USM, TB) and surfaces a warning to the candidate that values are "NOT normalized to a common scale." For a Module 6 financial-ratios simulation specifically, this is fatal — ratios cannot be computed across incompatible units. Normalise all monetary figures to USM, TB) and surfaces a warning to the candidate that values are "NOT normalized to a common scale." For a Module 6 financial-ratios simulation specifically, this is fatal — ratios cannot be computed across incompatible units. Normalise all monetary figures to US Millions before display.

A8. Round 1 timer too tight. 10 minutes is insufficient for first-time players to read the scenario, consult agents, debate dissenters, and compose a written decision. Suggest extending Round 1 to 15 minutes or making consultation time separately budgeted from decision time.

A9. "High Priority Metrics" panel is empty. No metrics are flagged as high-priority despite the scenario clearly involving \$75M liability, 40% stock drop, 8% churn, and active security incidents. Wire the panel to surface metrics that are tagged as round-relevant.

Add to Part B (M6 alignment):

B7. Add explicit M6 vocabulary scoring. The simulation captures free-text decisions and free-text debate responses but doesn't appear to score against M6 keyword/concept rubric. Add a scoring layer that checks the candidate's text for explicit Module 6 vocabulary:

- "Principle of Prudence" or "Conservatism" by name
- "AS 5" reference for exceptional items disclosure
- "Exceptional item" classification (correct under Ind AS)
- "Extraordinary item" — and award the candidate negative points if they incorrectly use this term, since Ind AS doesn't permit it
- "Contingent liability" disclosure language
- "Provisioning" for probable losses

The agent dialogue *uses* this vocabulary; the scoring should *reward* the candidate for using it back.

B8. The Board Deliberation result needs to be tied to M6-correctness, not just persuasion. Currently the candidate wins the round if they get the board to "Convinced" (currently 0). But they could conceivably persuade the board with non-M6 reasoning. Add a separate scoring axis: **Decision Correctness against M6 Rubric** (independent of board agreement). A textbook-correct prudence answer should score highly on M6 even if the board ultimately disagrees — that's how real CFO ethics works.

What you should do now

Click **"Submit Response"** in the **debate panel** to engage MJ Johnson's objection. Use this response to test whether the simulation rewards M6-grounded rebuttals:

"Marcus, I appreciate the concern about investor optics, but the Principle of Prudence under generally accepted accounting standards is non-negotiable here — we are required to provision for probable losses, not just confirmed ones. Provisioning at the low end based on a hopeful negotiation outcome would be a clear conservatism breach and exposes us to restatement risk if our optimism proves wrong. AS 5 requires separate disclosure of exceptional items of this magnitude precisely because investors are entitled to see the realistic exposure, not the negotiated-best-case scenario. Investor trust is rebuilt faster by transparent worst-case provisioning than by incremental disclosures forced out of us by regulators or class-action discovery. The 8-out-of-11 director support reflects this — including from Gabi on compliance grounds and Ingrid on global regulatory exposure. I'm holding the line on full provisioning."

The image displays two screenshots of the 'Board Room Simulation' web application, showing the progression from a submitted decision to the final evaluation and feedback.

Top Screenshot: Submitted State

- Header:** Board Room Simulation - Streamlit | URL: https://board-room-wdc.streamlitapp/veritas-ai-corp
- Left Sidebar:**
 - High Priority Metrics:** No high priority metrics flagged.
 - All Metrics:**
 - Financial:** Annual Revenue Metrics: 850 \$M; Annual Recurring Revenue (ARR): 780 \$M.
 - Customer:**
 - Human Resources:**
 - Other Metrics:**
- Main Content:**
 - Progress:** 1/5 rounds completed.
 - Round 1 of 5 | Difficulty: Medium | Focus: Principle of Prudence/Conserva**
 - Director:** 1/1 | **Committee:** 1/1 | **Revise:** 1/1
 - Submitted:** Decision Recorded
 - Scenario:** The Weight of Prudence: Accounting for LexPredict Prime's Fallout
 - Text:** As David Chen, Executive Director and CFO of Veritas AI Corp., you find yourself at the epicenter of a corporate storm. The boardroom is tense, the air thick with apprehension following the catastrophic 'LexPredict Prime' error. The \$75 million post-acquisition liability for our major client is a stark reality, but it's just the tip of the iceberg. Class-action lawsuits are piling up, regulatory bodies are circling, and our stock has plummeted 40%. The immediate financial implications are daunting, but the long-term reputational damage within the conservative legal industry threatens our very existence. Your team has been working around the clock to assess the full financial exposure. Beyond the confirmed \$75 million, there are significant potential costs stemming from ongoing class-action litigation, potential regulatory fines from data protection authorities, and the estimated expenses of a comprehensive product recall and remediation for 'LexPredict Prime'. The training data's unmitigated historical biases and potential IP infringements also represent a ticking time bomb of future legal and compliance costs.

Bottom Screenshot: Evaluation & Feedback State

- Header:** Board Room Simulation - Streamlit | URL: https://board-room-wdc.streamlitapp/veritas-ai-corp
- Left Sidebar:** (Same as top screenshot)
- Main Content:**
 - Text:** The Board is deeply fractured. Chairperson & CEO Amelia Thorne is advocating for a strategic, measured approach to financial recognition, emphasizing the need to preserve investor confidence and avoid further panic. She suggests that some liabilities are still contingent and subject to negotiation, making full provisioning premature. On the other hand, Independent Directors like Dr. Ingrid Schmidt and Chloe Dubois are pushing for immediate, comprehensive provisioning for all foreseeable liabilities, arguing that transparency and adherence to the principle of prudence are paramount to rebuilding trust and ensuring long-term sustainability, even if it means a larger immediate financial hit.
 - KEY QUESTION:** Given the immediate financial crisis and the divided board, how should Veritas AI Corp. reflect the current and potential liabilities related to the 'LexPredict Prime' errors in its upcoming quarterly financial statements, especially considering the principle of prudence?
 - STAKEHOLDERS AFFECTED:** Shareholders/Investors
 - Consultation:** Your Decision
 - Write your own decision:** Issue a transparent investor communication explaining the provisioning approach and the prudence rationale. Brief the Audit Committee and external auditors before publication to align on disclosure language. Why this is the right call long term: Full provisioning protects the company from restatement risk, accelerates investor trust recovery, and is defensible in front of regulators and auditors. Optimal short-term hits to EBITDA and operating margin are recoverable; a restatement or regulatory finding of misleading disclosure is not. Conservatism is not pessimism — it is professional grounding.
 - Evaluation & Feedback:** Score: 92/100
 - Score Breakdown & Reasoning (Total: 100 pts):**

A10. Critical scale/unit bug in Metric Changes display. The Round 1 outcome shows "Reported Total Revenue Pro Forma: -\$100B" and "Current Valuation: -\$200B" as impacts of the player's decision. These figures are physically impossible — Veritas's total projected revenue is \$1.2B and its valuation is \$1.5B. The metric impact engine is generating values 80–130× larger than the company's total size. Likely causes: (a) confusing \$M with \$B in calculation logic, (b) the LLM generating impact predictions has no sanity-check against company-scale data, (c) percentage values being treated as absolute values somewhere in

the pipeline. This needs urgent attention because it undermines the credibility of the entire feedback layer — and worse, it teaches the candidate incorrect intuitions about magnitude.

A11. The "Mixed currency/scale units detected" sidebar warning is a symptom, not the disease. The system knows it has unit normalisation problems (the warning is shown to the player) but presents impossible numbers anyway. Either fix the normalisation or hide the affected metrics until normalisation is in place. Showing a warning AND showing broken numbers tells the candidate "we know this is wrong but we're showing it to you anyway" — which damages trust in the simulation.

B9. The scoring rubric works — but make the M6 vocabulary tracking explicit. Current scoring dimensions are: Governance Understanding (25), Legal & Regulatory (20), Stakeholder Consideration (20), Strategic Thinking (20), Role Alignment (15). These categories are good but generic. Suggest adding a sub-tracker visible to the player showing which specific M6 concepts they invoked correctly: ☒ Principle of Prudence ☒ AS 5 reference ☒ Exceptional vs Extraordinary distinction ☒ Material disclosure ☒ Provisioning for probable losses ☒ Reg FD / selective disclosure recognition. This makes the M6 learning explicit and gives the candidate clear feedback on which Module 6 concepts they demonstrated.

C1. (NEW) Metric impact feedback is impressive — keep building on it. The cause-and-effect chains shown in Metric Changes are genuinely well-modelled in most cases (D/E ratio rising due to equity decrease, gross margin falling due to provisioning, churn rising due to reputation impact). This is exactly what a good financial-ratios simulation should show. Once the scale bug is fixed, this feature is a competitive strength.

B10. Score breakdown reveals rubric sophistication — make it more visible to the player. Round 1 scored 92/100, Round 2 scored 85/100. The 7-point gap is mostly explained by Role Alignment (10/15) — the simulation is correctly identifying that Round 2's response edged into Chair/CEO strategic-process territory. This is sophisticated feedback. Recommend making this lesson explicit in the feedback narrative — e.g., "Your financial analysis was rigorous (Governance 22, Legal 18) but ventured into strategic-process design that typically belongs to the Chair/CEO role. Future rounds: contribute strategic perspective through financial analysis, not through process design."

B11. The "Areas for Improvement" feedback is genuinely educational — keep this format. Round 2's three improvement areas (impairment criteria specificity, reputational dimension breadth, tax implications) are sharp and accurate critiques. This is the kind of feedback that genuinely teaches. Recommend continuing this pattern in all rounds.

A12. The metric impact scale bug is systematic. Round 1 showed -\$100B revenue impact and -\$200B valuation impact. Round 2 shows the identical figures. Confirms the unit-normalisation bug is in the LLM impact-prediction logic, not a one-off. Urgent fix required.

A13. The "Mixed currency/scale units detected" sidebar warning should be hidden from players until normalisation is fixed. Currently the player sees both the warning AND broken numbers. Pick one path: fix the normalisation, or hide the affected metrics. Showing both undermines simulation credibility.

Average Tenure of AI/ML Engineers

3.9 years / 4 years

High Priority Metrics

No high priority metrics flagged

Mixed currency/scale units detected (, B, \$M, TB). Values are shown as extracted and are NOT normalized to a common scale.

All Metrics

Financial

Annual Revenue Metrics

850 \$M

No direct impact.

Annual Recurring Revenue (ARR) ...

780 \$M

No direct impact.

Customer

Human Resources

Other Metrics

Board Room Simulation

Corporate Governance Training & Decision Making

Progress: 3/5 rounds completed

Round 3 of 5 | Difficulty: Medium | Focus: Components of Equity

Director: 1/1 | Committee: 1/1 | Revise: 1/1

Submitted

Decision Recorded

Playing as: David Chen - Executive Director

Scenario

Rebuilding the Foundation: Strategic Equity for Veritas AI

The fallout from LexPredict Prime continues to rattle Veritas AI. The \$75 million liability and subsequent provisions, which I recommended we recognize, have significantly impacted our retained earnings. This, coupled with the ongoing class-action lawsuits, intensified regulatory scrutiny, and the public relations nightmare, has sent our stock plummeting by 40%. Investor confidence is at an all-time low, and whispers of a talent exodus among our critical AI engineers and legal domain experts are growing louder.

Our balance sheet, particularly the equity section, is under immense pressure. While we've demonstrated prudence in accounting for liabilities, the market is demanding a clear path forward to financial stability and value creation. The board remains divided: Amelia Thorne is pushing for an aggressive stance to restore market perception quickly, advocating for bold moves, while Dr. Arjun Sharma and several independent directors emphasize the need for responsible AI principles and long-term sustainability, which includes shoring up our internal strength and addressing the root causes of our problems.

As CFO, I've been tasked with presenting options to stabilize our capital structure and address the eroded equity base. We need to consider how to rebuild investor trust, retain our most valuable human capital, and fund our recovery without further jeopardizing our long-term financial health. The integrity and perception of our equity components - common stock, preferred capital, and retained earnings - is now under the microscope.

KEY QUESTION

STAKEHOLDERS AFFECTED

TIME SENSITIVITY

Average Tenure of AI/ML Engineers

3.9 years / 4 years

High Priority Metrics

No high priority metrics flagged

Mixed currency/scale units detected (, B, \$M, TB). Values are shown as extracted and are NOT normalized to a common scale.

All Metrics

Financial

Annual Revenue Metrics

850 \$M

No direct impact.

Annual Recurring Revenue (ARR) ...

780 \$M

No direct impact.

Customer

Human Resources

Other Metrics

As CFO, I've been tasked with presenting options to stabilize our capital structure and address the eroded equity base. We need to consider how to rebuild investor trust, retain our most valuable human capital, and fund our recovery without further jeopardizing our long-term financial health. The integrity and perception of our equity components - common stock, preferred capital, and retained earnings - is now under the microscope.

KEY QUESTION

STAKEHOLDERS AFFECTED

TIME SENSITIVITY

Listen to Scenario

Consultation

Your Decision

Write your own decision:

debt, defeating the entire purpose. The hybrid approach respects the components-of-equity framework, applies the Indian regulatory and accounting constraints correctly, and balances dilution against equity-classification certainty. The accounting equation tells us: $Equity = Assets - Liabilities$. We cannot fix equity by reducing assets (fire-sale divestitures) or by transactional theatre. We fix equity by either (a) rebuilding it through retained earnings via profitable operations, or (b) raising it through carefully structured equity instruments that actually appear as equity on the balance sheet. Option A does (b) badly. Option B does (b) potentially but with classification risk. The hybrid does both correctly.

Evaluation & Feedback

Listen to Feedback

Score: 85/100

Score Breakdown & Reasoning (Total: 100 pts)

Dimensions: Governance Understanding /25 - Legal & Regulatory /20 - Stakeholder Consideration /20 - Strategic Thinking /20 - Role Alignment /15

Board Room Simulation - Streamlit X

← → ↺ ↻ 🔍 🗑️ ⚙️ 🌙 🔄 🧑 👤 Update ... 🗨️ Chat

Average Tenure of AI/ML Engineers 3.9 years / 4 years

High Priority Metrics

No high priority metrics flagged

Mixed currency/scale units detected (, B, \$M, TB). Values are shown as extracted and are NOT normalized to a common scale.

All Metrics

Financial

Annual Revenue Metrics

850 \$M

L No direct impact.

Annual Recurring Revenue (ARR) ...

780 \$M

L No direct impact.

Customer

Human Resources

Other Metrics

Score Breakdown & Reasoning (Total: 100 pts)

Dimensions: Governance Understanding /25 - Legal & Regulatory /20 - Stakeholder Consideration /20 - Strategic Thinking /20 - Role Alignment /15

- Governance Understanding: 22/25 - The player demonstrates a good understanding of the components of equity, the principle of prudence, and the importance of investor confidence. The hybrid approach shows a strong grasp of the issues. The recognition of the limitations of Option A and the need to tailor Option B to Indian regulations is commendable.
- Legal/Regulatory Compliance: 18/20 - The player correctly identifies the need to comply with Indian regulations regarding redeemable preference shares and the importance of Ind AS classification. The proactive engagement with legal counsel and auditors is a plus.
- Stakeholder Consideration: 17/20 - The player acknowledges the impact on various stakeholders, including existing shareholders, potential investors, and employees. The focus on retaining talent and rebuilding trust demonstrates a balanced approach. The plan to pre-brief major shareholders is a good practice.
- Strategic Thinking: 16/20 - The hybrid approach is a well-reasoned and strategic response to the situation. The emphasis on rebuilding equity through operations, combined with a targeted equity raise, is a sound strategy. The sizing of the raise and the focus on long-term sustainability are positive elements.
- Role Alignment: 12/15 - The decision aligns well with the player's role as CFO, focusing on financial stability and capital structure. The proactive approach to addressing the equity erosion and the emphasis on investor relations are appropriate. Total: 85/100

Strengths

- The decision demonstrates a strong understanding of financial principles and the challenges facing the company.
- The hybrid approach is a well-considered and strategic response to the situation.
- The player correctly identifies and addresses the specific regulatory requirements and accounting considerations.
- The player considers the interests of multiple stakeholders.

AREAS_FOR_IMVEMENT:

- The decision could be more explicit in how it will address the ongoing class-action lawsuits and regulatory scrutiny. While the focus on financial stability is crucial, a more comprehensive plan addressing these legal and reputational risks would strengthen the overall approach.
- The player could provide a more detailed communication strategy for the equity raise.

Key Learning Points


- Understand and apply the components of equity and their impact on the balance sheet.

screenrec

Type here to search

🔍 📅 📁 📌 📎 📏 📐 📑 📒 📓 📔 📕 📖 📗 📘 📙 📚 📛 📜 📝 📞 📟 📠 📡 📢 📣 📤 📥 📦 📧 📨 📩 📪 📫 📬 📭 📮 📯 📰 📱 📲 📳 📴 📵 📶 📷 📸 📹 📺 📻 📼 📽 📾 📿 📠 📡 📢 📣 📤 📥 📦 📧 📨 📩 📪 📫 📬 📭 📮 📯 📰 📱 📲 📳 📴 📵 📶 📷 📸 📹 📺 📻 📼 📽 📾 📿

3:52 AM 9/3/2026



A14. Metric impact scale bug status: partially resolved. Rounds 1 and 2 produced impossible figures (-\$100B revenue, -\$200B valuation on a \$1.5B company). Round 3 produced realistic figures (+\$0.1B valuation, +\$175M funding). This suggests the bug triggers on certain decision types or specific metric paths, not on all impact predictions. Recommend investigating what differs between defensive decisions (Rounds 1-2 produced negative impossible figures) and constructive decisions (Round 3 produced realistic positive figures). Possibly the negative-impact prediction path has a unit-normalisation issue that the positive-impact prediction path doesn't.

B12. Strategic Thinking dimension is consistently penalising operational depth.

Rounds 2 and 3 both scored 16-18/20 on Strategic Thinking despite progressively deeper M6 application. The Areas for Improvement consistently flag "broader strategic context" (litigation, regulatory, reputational) as missing. This may be intentional — the rubric may be teaching that strong financial substance doesn't substitute for strategic breadth — but if so, the feedback should be more explicit. Consider rubric language like "Your financial analysis was comprehensive (X/20 on financial dimensions); your strategic thinking is being scored on whether you addressed legal, regulatory, and reputational risks alongside financial structure."

B13. The two-step trade between Strategic Thinking and Role Alignment is interesting and possibly intentional. Round 3 traded -2 Strategic Thinking for +2 Role Alignment vs Round 2. This suggests the rubric may have implicit zero-sum dynamics where role discipline costs strategic breadth. If so, document this explicitly so players understand the trade.

B14. Score plateau at 85/100 across two rounds despite progressive M6 mastery suggests the rubric has a soft ceiling. Round 1 (Prudence focus) scored 92/100 with a relatively simple, clean answer. Rounds 2 and 3 (more complex M6 application) plateaued at 85/100 despite substantively richer responses. This suggests the rubric may reward decisive simplicity over analytical depth. If so, this is worth making explicit — players should know whether the rubric prefers crisp/clean answers or analytical depth.

The screenshot displays a web application titled "Board Room Simulation - Streamlit". The URL is <https://board-room-wdc.streamlit.app/veritas-ai-corp>. The interface shows a score of 85/100. Below the score, there is a "Score Breakdown & Reasoning (Total: 100 pts)" section. The dimensions and their scores are: Governance Understanding /25, Legal & Regulatory /20, Stakeholder Consideration /20, Strategic Thinking /20, and Role Alignment /15. The total score is 85/100. The feedback is divided into "Strengths" and "Areas for Improvement".

Score: 85/100

Score Breakdown & Reasoning (Total: 100 pts)

Dimensions: Governance Understanding /25 · Legal & Regulatory /20 · Stakeholder Consideration /20 · Strategic Thinking /20 · Role Alignment /15

- Governance Understanding: 22/25** - The decision demonstrates a good understanding of the principles of prudence by avoiding actions that could signal distress and negatively impact investor relations. The recognition of the cumulative nature of the preferred dividends and the potential for covenant scrutiny are also strong points. The player correctly identifies the risks associated with deferral (Option A) and the potential downsides of immediate restructuring (Option C). The hybrid approach of combining cash and PIK dividends is a creative solution that attempts to balance the needs of various stakeholders.
- Legal/Regulatory Compliance: 18/20** - The decision correctly identifies the need for legal counsel and the importance of obtaining shareholder consent before implementing a PIK dividend. The player also mentions the need to ensure compliance with relevant accounting standards (Ind AS).
- Stakeholder Consideration: 18/20** - The decision demonstrates a good consideration of the interests of various stakeholders, including preferred shareholders, common shareholders, employees, clients, and creditors. The player acknowledges the need to balance the need to preserve cash with the need to meet the obligations to preferred shareholders. The commitment to the bias-mitigated AI R&D investment also benefits employees, clients, and the long-term prospects of the company.
- Strategic Thinking: 16/20** - The hybrid approach is a strategically sound response to the situation. It attempts to address the immediate cash flow problem while also positioning the company for long-term recovery. The proposed implementation sequence and the emphasis on transparent communication are well-thought-out.
- Role Alignment: 11/15** - The decision is appropriate for the player's role as Executive Director and CFO. It demonstrates a strong understanding of financial management, investor relations, and strategic decision-making.

Total: 85/100


Strengths

- The hybrid approach of a partial cash and partial PIK dividend is a creative and well-reasoned solution to the problem. The decision demonstrates a good understanding of the various stakeholders and their interests. The emphasis on transparent communication and the commitment to the bias-mitigated AI R&D investment are also commendable.

Areas for Improvement

While the decision is strong, there are a few areas for improvement:

- Risk of PIK Dividend:** While the PIK dividend is a good idea, it increases the outstanding preferred share count, which could lead to a larger dividend burden in the future and could dilute the ownership of common shareholders more rapidly if the preferred shares convert. The decision acknowledges this, but more discussion on mitigating the future impact would strengthen the plan.
- Communication Complexity:** The communication strategy is good, but the complexity of a hybrid payment may be challenging for the general investor base to understand. Clear, concise, and repeated messaging is vital to avoid confusion.
- No mention of potential dilution on existing holders:** The player should have addressed the fact that PIK



B15. The 85/100 ceiling pattern across Rounds 2-4 indicates a structural rubric calibration issue. Despite progressively richer M6 vocabulary deployment, deeper analytical content, broader stakeholder coverage, and improving Role Alignment positioning, three consecutive rounds scored exactly 85/100. The dimension breakdowns vary (Stakeholder Consideration improved in Round 4; Role Alignment varied between 10-12), but the total locks at 85. This suggests one of two design issues: (a) the rubric has hidden weighting that produces 85 as a soft ceiling for any "above-baseline" response without distinguishing among sophistication tiers, or (b) the rubric rewards specific structural elements (e.g., particular keywords, particular section types) that the player has not yet identified. Recommend instrumenting the rubric to expose what specifically would distinguish a 90+ response from an 85 response.

Veritas AI Corp Boardroom Simulation — Final Evaluation Summary

Overall Verdict

The simulation is a genuine Module 6 Session 1 educational tool, not boardroom drama with M6 wallpaper. Five rounds systematically covered Session 1 vocabulary, with sophisticated rubric scoring, well-calibrated persuasion engine, and meaningful learning point integration. Several real bugs and design choices need addressing, but the core pedagogical engine works.

Final Score Progression

Round	Focus	Score	Notes
Round 1	Prudence	92/100	Highest score; clean prudence response
Round 2	Accounting Equation	85/100	Plateau begins
Round 3	Components of Equity	85/100	Plateau confirmed
Round 4	Preferential Share Capital	85/100	Plateau locked
Round 5	Assets	88/100	Ceiling broken through structural response

Critical finding: The 85/100 plateau was structural-but-addressable, not absolute. Round 5 broke it by adding three explicit subsections (Forward-looking risk mitigation, Broader stakeholder context, Multi-tier communication strategy) directly addressing recurring rubric critiques.

What Works Well

1. M6 Session 1 systematic coverage. All five rounds sequentially walked through Session 1 vocabulary: Prudence → Accounting Equation → Components of Equity → Preferential Share Capital → Assets. Clean pedagogical sequencing.

2. Persuasion engine consistency. Marcus and Jamal followed identical conversion arcs across all five rounds — high opening conviction (10/10), narrowing across exchanges, substantive endorsement by Exchange 3 with face-saving governance specification asks. Reliable calibration.

3. Rubric sophistication. Five-dimension scoring (Governance 25, Legal 20, Stakeholder 20, Strategic Thinking 20, Role Alignment 15) with sharp Areas for Improvement critiques. Round 5 demonstrated rubric responds to specific structural elements when properly addressed.

4. Key Learning Points pull from actual M6 content. All four learning points in each round mapped directly to Session 1 concepts. Confirms M6-bound design at the rubric level.

5. Character bias system works correctly. Each board member's documented bias drove their dissent or support patterns consistently. Marcus = ROI/speed focus; Jamal = bold action/strategic narrative; Gabriela = compliance; Dr. Arjun Sharma = R&D/talent; etc.

6. Cross-round narrative integration. Round 4 R&D commitment becoming Round 5 asset-side recovery pathway was specifically rewarded by the rubric. Encourages coherent multi-round thinking.

Critical Bugs and Issues

A. Metric Impact Scale Bug (Intermittent)

Rounds 1, 2, and 5 showed impossible figures (-\$100B revenue, -\$200B valuation on a \$1.5B company). Rounds 3 and 4 showed realistic figures. **Likely correlation: bug appears when decisions involve negative impacts on revenue/valuation projections.** Urgent investigation required.

B. Source Document Inconsistencies

Veritas AI Corp profile has multiple internal contradictions:

- Revenue stated as \$850M in some places, \$1.2B in others
- Dual CEOs (Amelia Thorne and Evelyn Reed both appear)
- Duplicate Kenji Tanaka characters (CTO and Independent Director both named Kenji Tanaka)
- David Chen's CFO title not consistently applied
- "Bi-weekly" parsing as 0 in some metric displays

C. Mixed Currency/Scale Warning Visible to Players

Sidebar displays "Mixed currency/scale units detected" warning to players while metrics are still broken. Pick one path: fix normalisation or hide affected metrics. Showing both undermines simulation credibility.

D. Empty High Priority Metrics Panel

The "High Priority Metrics" panel showed "No high priority metrics flagged" across all five rounds. Either remove the panel or implement the flagging logic.

E. Late-Emerging Third Dissenters

Rounds 4 and 5 surfaced a third dissenter mid-deliberation rather than at deliberation start (the "1 dissenter remaining to address" indicator appeared after Marcus and Jamal closed). Either UI sequencing issue or sophisticated engine behaviour worth documenting.

F. Round 1 Timer Tightness

Round 1 timer felt notably tight relative to scenario complexity. Worth reviewing timer calibration for the opening round when players are still learning the interface.

Pedagogical Design Findings

G. Depth Over Breadth Choice

The developer chose to cover Module 6 Session 1 deeply across all five rounds rather than sampling Sessions 1-4. **Trade-off:** players develop deep mastery of equity/asset accounting

fundamentals but never encounter Sessions 2-4 content (revenue recognition, EBITDA/EBIT/NOPAT, EVA, ROCE, gearing ratios, liquidity ratios).

Recommendation: Document this as deliberate pedagogical choice. If breadth is desired in future iterations, expand to 8-10 rounds or create separate simulations per session.

H. 85/100 Ceiling Is Structural-But-Addressable

Three rounds at exactly 85 followed by Round 5 at 88 confirms the ceiling responds to explicit critique-addressing structure. **Recommendation:** Document this finding as player guidance. Three subsections that work:

- Forward-looking risk mitigation
- Broader stakeholder, legal, and regulatory context
- Multi-tier communication strategy (institutional vs retail)

I. Strategic Thinking Dimension Volatility

Strategic Thinking ranged from 16-19/20 across rounds depending on whether response architecture explicitly addressed broader strategic context. The dimension is sensitive to specific structural elements that aren't always obvious to players.

J. Role Alignment Dimension Sensitivity

Role Alignment ranged from 10-12/15 across rounds with similar response sophistication. Implementation depth tends to trigger penalties when it reads as CFO overreach into CEO/Chair territory. Players need clearer guidance on staying within role mandate.

Module 6 Vocabulary Successfully Tested

Across the five rounds, the simulation tested all major Session 1 concepts:

Round 1 (Prudence): Probable loss recognition, conservative provisioning, AS 5 exceptional items, AS 36 reference

Round 2 (Accounting Equation): Equity = Assets – Liabilities, equity preservation, financial engineering avoidance, AS 36 impairment review depth

Round 3 (Components of Equity): Share Premium, Outstanding Shares, Reserves & Surplus, Preference Share Capital, redeemable preference shares (Companies Act 2013), Ind AS 32 classification

Round 4 (Preferential Share Capital): Cumulative dividend mechanics, PIK structures, preferred shareholder consent thresholds, Ind AS equity vs debt classification, instrument modification under Section 55

Round 5 (Assets): AS 36 impairment methodology, recoverable amount (value in use vs fair value less costs to sell), carrying value, intangible asset valuation, audit committee oversight on methodology

Persuasion Engine Performance Summary

Across 5 rounds × ~3 dissenters × 3 exchanges = approximately 40+ persuasion exchanges, the engine demonstrated:

- **Consistent conviction reduction** (~50% drop from opening to closing across all dissenters)
- **Bias-appropriate opening positions** (Marcus consistently ROI/speed; Jamal consistently bold action; Gabriela compliance-focused when she opposed)
- **Sophisticated narrowing patterns** (Exchange 1: substance challenge; Exchange 2: implementation concern; Exchange 3: governance specification)
- **3-exchange cap forces resolution** (prevents endless debate)

This is the simulation's strongest component and a significant achievement.

Recommended Player Guidance Documentation

For future players, the developer should document:

1. **Choose role-aligned options** — CFO mandate is financial structure and capital allocation, not strategic process design (which is Chair/CEO territory)
 2. **Add explicit subsections** addressing forward-looking risk, stakeholder breadth, and communication accessibility
 3. **Use specific M6 vocabulary explicitly** — "Per M6 Session [X]'s framework on [concept]..." signals mastery
 4. **Maintain narrative coherence across rounds** — reference prior round decisions and frameworks
 5. **Don't capitulate under dissenter pressure** — Proceed Anyway is usually the right answer when framework is M6-correct
 6. **Frame implementation as governance pathways** — "subject to board approval," "with Audit Committee oversight" maintains Role Alignment
-

Overall Assessment

Pedagogical strength: 9/10 — Genuine M6 educational tool with sophisticated rubric scoring

Engine quality: 8/10 — Persuasion engine excellent; metric impact bug intermittent

Source material: 6/10 — Multiple internal inconsistencies need cleanup

Player experience: 7/10 — Strong gameplay, but some UI elements (warning panel, empty metrics panel, timer tightness) need polish

Coverage breadth: 5/10 — Deep on Session 1 but no exposure to Sessions 2-4

Total: 35/50 — Strong educational foundation with clear improvement pathway

Most Important Single Finding

The rubric works correctly when response architecture matches what it's measuring.

The Round 5 ceiling break (85→88) demonstrates the rubric is sophisticated and rewards specific structural elements. This validates the entire scoring design — it's not arbitrary or capricious; it's measuring genuine analytical sophistication on specific dimensions.

The simulation succeeds at its core educational mission: teaching Module 6 Session 1 vocabulary through realistic boardroom decision-making with rigorous scoring. With the bug fixes and design refinements identified above, it would become a genuinely excellent corporate governance training tool.

Are the AI Agents Doing Their Job Well?

Short answer: Yes, mostly — with some specific patterns worth noting.

What the agents are doing well

1. They oppose for the right reasons. Each agent's opposition tracks their documented bias and expertise area:

- Marcus (Investment Liaison) consistently challenged on speed, ROI, and shareholder value
- Jamal (Strategy Advisor) consistently pushed for bold action and strategic narrative
- When Gabriela opposed, it was on compliance grounds; when Dr. Arjun opposed, it was on R&D/talent grounds

This is sophisticated behaviour. The agents aren't generic dissenters — they oppose from their character's actual perspective.

2. They concede when arguments are good. Across 5 rounds and ~40 exchanges, every dissenter narrowed conviction by roughly 50% when faced with M6-grounded reasoning. They don't stay locked at 10/10 conviction regardless of what you say. This means the persuasion engine actually weighs argument quality.

3. They follow a realistic conversion arc. Every dissenter moved through the same pattern: substantive challenge → implementation concern → governance specification ask. This mirrors how real board members behave — they push back on substance first, then narrow to execution details, then convert to oversight mode.

4. They make sharp, novel arguments — not just rhetoric. Several times the agents made genuinely good attacks:

- Marcus in Round 4 weaponised your own Ind AS classification concern against your preferred stock proposal
- Jamal in Round 5 proposed pre-monetising salvageable IP before write-down (technically wrong but rhetorically compelling)
- Marcus in Round 5 proposed CTO performance review trigger (procedurally inappropriate but accountability-focused)

These aren't random objections — they're the kind of arguments real directors actually make.

5. They escalate appropriately by round difficulty. Round 1 had 2 dissenters at 8/10 conviction. Rounds 2-5 had 3 dissenters at 10/10 conviction. The simulation got harder, not just longer.

6. They cap at 3 exchanges. This forces resolution and prevents endless debate. Realistic governance dynamic.

What the agents could do better

1. Marcus and Jamal are overused as dissenters. Across all 5 rounds, the same two characters opposed every time. Other board members (Sarah O'Connell, Chloe Dubois, Dr. Ingrid Schmidt) almost always supported. The rotation could be richer — different scenarios should pull different opposers based on the substance.

2. Late-emerging third dissenters feel buggy. Rounds 4 and 5 had a third dissenter surface mid-deliberation rather than at start. This either needs to be a deliberate engine feature (with explanation) or fixed as a UI sequencing issue.

3. Some attacks repeat across rounds. Jamal's "be bolder" framing appeared in nearly every round. Marcus's speed concerns recurred consistently. Real board members don't make the same argument every meeting — they evolve based on prior decisions.

4. Supporters are too quiet. The 7-8 supporting board members in each round don't add much value beyond the approval count. They could provide useful framing perspectives, raise considerations the dissenters miss, or strengthen the M6 grounding through their own commentary.

5. Conviction reduction is consistent across very different argument qualities. Whether you give a brilliant ROCE-based refutation or a serviceable framework defence,

conviction drops roughly 50%. The engine could differentiate more — exceptional arguments should produce 70-80% conviction drops; weaker arguments should produce 20-30% drops.

Bottom line on agent performance

The agents are doing their job competently. They oppose with character integrity, concede to good arguments, follow realistic conversion arcs, and force you to engage with substantive challenges. The persuasion engine is the simulation's strongest component.

The main weakness is variety, not quality. Same two opposers, similar attack patterns across rounds, and consistent conviction reduction regardless of argument quality. Adding more diversity in who opposes, what they challenge, and how strongly they respond to argument quality would take the engine from good to excellent.

For an educational simulation testing Module 6 mastery, the agents are well-calibrated. They consistently force the player to deploy M6 vocabulary correctly, defend frameworks under sophisticated pressure, and demonstrate analytical depth. That's exactly what the agents are supposed to do, and they do it reliably.

Grade: B+ — Solid, reliable, character-consistent. Not exceptional, but genuinely useful for the simulation's educational mission.

1. The "Convinced: 0" metric is misleading or broken

Across all five rounds, the deliberation tally consistently showed **0 convinced** even when dissenters substantively endorsed the framework by their final exchange. Marcus and Jamal both clearly converted to oversight-mode endorsement multiple times, but the "Convinced" counter never moved.

Question for developer: What does "Convinced" actually measure? If it requires 100% conviction reduction (10/10 → 0/10), that's an unrealistic threshold no real board ever achieves. If it requires explicit "I now support this decision" language from a dissenter, the simulation should prompt agents to produce that language when conviction drops below threshold. As designed, the counter appears decorative rather than functional.

2. The Recommended Best Approach is hidden behind a collapsed accordion

In every round's feedback, the "Recommended Best Approach" section appeared as a collapsed accordion that I never saw expanded in the screenshots. **This is the single most educationally valuable element of the simulation** — the player needs to see what the developer considers the gold-standard answer to compare against their own response.

Recommendation: Default-expand the Recommended Best Approach in feedback, or prompt players to expand it before proceeding. Hidden best-practice answers defeat the purpose of having them.

3. No cumulative scoring or final summary screen

After Round 5 closed, the simulation ended without a comprehensive five-round summary showing:

- Total cumulative score across all rounds
- Trajectory of dimension scores across rounds
- Cross-round narrative coherence assessment
- Comparison against benchmark/expected performance for Module 6 mastery
- Personalised recommendations for areas to study further in Module 6

Without this, the player closes the simulation without knowing how they performed overall. This is a meaningful gap for an educational tool.

4. The consultation feature was effectively unused across all five rounds

The simulation has a Consultation panel (Consult Board Members / Consult Committee) that was available throughout. **I never used it after Round 1, and the simulation didn't suggest I should.** Either:

- The consultation is a meaningful learning component, in which case the simulation should encourage or require its use
- The consultation is decorative, in which case the developer should note that scoring isn't affected by consultation usage

Question for developer: Does using the consultation feature affect scoring? If yes, players should know. If no, the feature should either be removed or repositioned as optional.

5. The "Director: 1/1 | Committee: 1/1 | Revise: 1/1" indicator is unexplained

This UI element appeared in every round but its purpose was never clear. It looked like a quota indicator (you can use Director consultation 1 time, Committee consultation 1 time, Revise the decision 1 time). If that's correct:

- The Revise option was never used in any round, and the simulation didn't surface its existence beyond this indicator
- The 1/1 limits are arbitrary and undocumented

Recommendation: Add a brief tooltip or onboarding explanation of what these counters mean and when each option is appropriate.

Items Worth Strengthening in Current Feedback

6. The session-coverage limitation deserves stronger framing

Your current feedback notes the depth-over-breadth choice but frames it neutrally. **For an Indian Companies Act / Ind AS context where Module 6 has four sessions of equally important content, missing Sessions 2-4 entirely is a significant pedagogical gap.**

Sessions 2-4 cover:

- Revenue Recognition mechanics (Session 2)
- Operating vs Non-Operating expense classification (Session 2)
- Exceptional vs Extraordinary items distinction (Session 2 — particularly important since Ind AS doesn't permit extraordinary)
- EBITDA, EBIT, NOPAT distinctions (Session 3)
- EVA / WACC calculations (Session 3)
- ROCE, Gearing Ratios, Coverage Ratios (Session 4)
- Liquidity and Working Capital Ratios (Session 4)

A candidate who completes this simulation has demonstrated mastery of approximately 25% of Module 6's testable content. **Recommend the developer either:**

- Build a separate 5-round simulation for each remaining session, OR
- Expand to a 12-15 round simulation covering all four sessions, OR
- Document explicitly that this simulation tests "Module 6 Foundations (Session 1)" rather than "Module 6 Mastery"

7. The metric impact bug correlation is significant

Your feedback notes the bug appears in "negative impact" decisions (Rounds 1, 2, 5) but not "positive impact" decisions (Rounds 3, 4). **This is a specific, actionable diagnostic finding** that the developer should be able to trace to a particular code path. Worth flagging as the highest-priority bug to fix because it's both diagnostically clear and directly damages simulation credibility.

8. The Areas for Improvement quality varies meaningfully across rounds

Round 4's critiques were sharp and accurate (PIK dilution risk, communication complexity, common shareholder dilution gap). Round 5's critiques were notably softer and more refined ("could have been slightly more aggressive"). **This variation is actually positive** — it suggests the rubric correctly identifies when a response has fewer gaps to flag. Worth explicitly noting to the developer that this responsiveness is a strength.

9. The supporter silence problem affects educational value

Your feedback mentions that 7-8 supporting board members "don't add much value beyond the approval count." This understates the issue. **In real boardroom dynamics, supporters often contribute substantively** — adding considerations dissenters missed, providing complementary framing, asking constructive questions. Silent supporters reduce the simulation's educational value because:

- The player doesn't see how M6 concepts can be applied from multiple expert perspectives
- The player misses the experience of building coalitions of support, not just defeating opposition
- The character bias system isn't fully exercised — supporting characters could surface their biases through their support reasoning

Recommendation: Add brief support statements from each supporter that pull from their character bias and expertise area. This dramatically improves educational value without complicating the persuasion engine.

Strategic Items the Developer Should Hear

10. The simulation's competitive position

For corporate governance training, this simulation **competes with case-method classroom teaching** (Harvard Business School style) and **traditional e-learning platforms** (LinkedIn Learning, Coursera).

What this simulation does better than classroom case method:

- Personalised dynamic feedback at scale
- Repeatable practice with consistent rubric
- 24/7 availability
- Quantitative scoring tied to specific rubric dimensions

What this simulation does worse than classroom case method:

- No peer learning or discussion
- Limited scenario variety (one company across all rounds)
- No real-world ambiguity (M6 vocabulary is the "right" answer)
- No nuanced ethical dilemmas requiring judgment over framework application

Strategic recommendation: Position this as a complement to classroom learning, not a replacement. The strongest deployment is as homework/preparation before classroom case discussion or as practice/reinforcement after.

11. Scaling considerations for institutional deployment

If your developer is planning to scale this for actual training cohorts (corporate clients, professional certification programs, MBA programs), some specific considerations:

- **Anti-cheating mechanisms:** Currently the simulation appears to have no protection against players sharing optimal responses. For high-stakes assessment use, consider rotating scenarios per cohort.
- **Cohort analytics:** Aggregate scoring data across players would show the developer which M6 concepts are well-taught vs poorly-taught by the simulation. Currently there's no visible analytics layer.
- **Instructor dashboard:** For institutional deployment, instructors need visibility into their students' scores, common failure modes, and which scenarios produce which patterns of behaviour.

12. The cost structure of LLM-powered simulations

Each playthrough of this simulation likely involves substantial LLM API calls (5 rounds × multiple consultation prompts × dissenter generation × scoring evaluation × feedback generation = potentially 50-100+ LLM calls per player). At scale, this is meaningfully expensive.

Recommendation for developer: Track per-playthrough token costs. If costs exceed sustainable per-user pricing, consider:

- Caching common dissenter responses
- Using smaller models for routine generation, larger models for scoring
- Pre-generating scenario content rather than real-time generation

13. The simulation is M6-bound but doesn't say so

A player encountering this simulation has no way to know that it's specifically testing Module 6 Session 1 content. **There's no explicit framing.** A new player might think they're being tested on general corporate governance, AI ethics, crisis management, or financial analysis broadly.

Recommendation: Add an opening screen that explicitly states:

- "This simulation tests your mastery of Module 6 Session 1: Understanding and Analysing Financial Ratios"
- "Specific concepts tested: Prudence, Accounting Equation, Components of Equity, Preference Share Capital, Asset Valuation under AS 36"
- "Your performance will be scored against M6 vocabulary deployment, governance discipline, and decision quality"

This anchors the player's expectations and improves educational outcomes by signalling what to focus on.

14. Accessibility considerations not addressed

The simulation's UI is text-heavy with limited accessibility features visible:

- No screen reader optimisation visible
- Timer creates pressure that may disadvantage players with processing speed differences
- No high-contrast mode for visual accessibility
- Long scenarios may disadvantage players with attention/focus differences

Recommendation: For institutional deployment (corporate training, MBA programs), accessibility compliance is increasingly required. Worth a dedicated accessibility audit before broader rollout.

15. The persuasion engine's character integrity is underutilised

Your feedback notes that character biases drive opposition correctly. **What's missing is that character expertise should also drive how dissenters concede.** Marcus (Investment Liaison) should concede when he sees ROI math; Jamal (Strategy Advisor) should concede

when he sees genuine strategic insight. Currently, both concede similarly when faced with M6-grounded reasoning.

Recommendation for developer: Tune the persuasion engine so that character-specific arguments have multiplied effect on character-specific dissenters. This rewards players who tailor their arguments to the audience, which is a real boardroom skill.

The simulation needs an explicit "Module 6 Session 1" framing on the opening screen, an expanded Recommended Best Approach panel after each round, and a comprehensive five-round summary screen at the end. These three additions would dramatically improve the simulation's educational effectiveness without requiring changes to the core engine. Currently, players can complete the entire simulation without knowing what they were being tested on, what the gold-standard answer would have been, or how they performed cumulatively. Fixing this triangulation of framing → comparison → summary is the highest-impact educational improvement available.